

Lecture 11

Relational Algebra (cont'd)

Week 5

Overview

- Join
 - join
 - equijoin and natural join
- Division
- Simplifying relational algebra expressions

Join

$$A \bowtie_C B$$

Not strictly necessary, but very useful

$$A \bowtie_C B \equiv \sigma_C (A \times B)$$

I.e. take the cartesian product (cross product) of A and B and select tuples which satisfy condition C

Equijoin

$$A \bowtie_{\text{equality conditions}} B$$

where the equality conditions prescribe that corresponding attribute values be equal

E.g. **Employee** $\bowtie_{\text{SSN = ESSN}}$ **Dependent**
will join the corresponding tuples of
Employee and Dependent

ssn	essn

A

B

Unnecessary duplicates

↓

	ssn	essn	
	1	1	
	2	2	
	3	3	
		

A

▷◁

ssn=essn

B

Natural join

- To remove the duplicate columns a variation of equijoin is used

$A *_{\text{equality conditions}} B$

Example:

$\text{Employee} *_{\text{ssn=essn}} \text{Dependent}$

- If no condition is listed, *all corresponding attributes* (ones with the same name) are joined through the equality condition
But, be careful:

Employee * Dependent

- a) would not work, because there are no attributes with the same name
- b) sometimes attribute names inadvertently join. *It is best to always list the join condition explicitly!*

Division

$$A \div B$$

Example:

“List the employees who work on every project”

Take a simplified company schema:

employee(ssn)

project(pno)

works_on(ssn,pno)

- If we knew that every employee works on every project, then the works_on relation would be unnecessary, it could be reconstructed by the cartesian product

$$W = \text{employee} \times \text{project}$$

- if we took away works_on from W:

$$W \setminus \text{works_on}$$

the result would contain potential, but not actual works_on relationship instances


- Therefore


$$\Pi_{ssn} (\text{Employee} \times \text{Project} \setminus \text{Works_on})$$


is the set of employees who do *not* work on every project

- To calculate the set of employees who work on every project we take this away from employee:


$$\Pi_{ssn} \text{Employee} \setminus \Pi_{ssn} (\text{Employee} \times \text{Project} \setminus \text{Works_on})$$







Employees who do *not* work on every Project



Employees who work on every Project

Introduce the division operator:

$$\begin{aligned} \text{Works_on}(\text{ssn}, \text{pno}) \div \text{Project}(\text{pno}) \equiv \\ \Pi_{\text{ssn}} \text{Employee} \setminus \\ (\Pi_{\text{ssn}} (\text{Employee} \times \text{Project} \setminus \text{Works_on})) \end{aligned}$$

We can do this in general as well:

Definition

$$\begin{aligned} A(a,b) \div B(b) \equiv \text{the largest set } C(a) \\ \text{such that } B \times C \subseteq A \end{aligned}$$

Simplifying relational algebra expressions

- $10 * 5 + 10 * 7 + 10 * 8 =$
 $10 * (5 + 7 + 8) = 10 * 20 = 200$
- In the same way expressions in relational algebra can be simplified based on the discussed properties of the operators
- Query processors evaluate queries using simplification or to the contrary, elaboration so that the execution is cheaper (e.g cutting down on tuples which will later be discarded anyway)

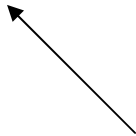
Example

Simpler



$\Pi_{ssn,name} ($
 $\sigma_{name='mary', Employee} *_{dno=dnumber}$
 $\sigma_{dname='research', Department}) =$

$\Pi_{ssn,name} ($
 $\Pi_{ssn,name,dno} \sigma_{name='mary', Employee} *_{dno=dnumber}$
 $\Pi_{dnumber} \sigma_{dname='research', Department})$



Cheaper to evaluate

Query plan

- Relational algebra operators are implemented as parametric programming language functions, manipulating data structures that implement sets of tuples (relations)
- A query plan is a suitable concatenation (or parallel execution model) of such functions

Query plans can be evaluated ..

- ...by the main processor of the computer on which the DBMS runs
- Simple operations, like select or project, can be performed by the disk drive itself - in order to avoid the transfer of massive amounts of data to main memory (later to be discarded anyway)

Conclusion

- Relational algebra has been introduced as an algebra of relations
- Each operator (monadic or binary) works on relations and returns relations
- Additional operations exist to carry out aggregation (this will be discussed only when studying SQL).
- Relational algebra can be used to express queries, as expressions to be evaluated on a database instance
- **Caution:** The result will depend on the actual database instance. The expression should be correct for *every* / *any* database instance.

The end