

Lecture 11

Relational Algebra (cont'd)

Week 5

Overview

- Join
 - join
 - equijoin and natural join
- Division
- Simplifying relational algebra expressions

© 2010 Griffith University

2

Join

$$A \bowtie_C B$$

Not strictly necessary, but very useful

$$A \bowtie_C B \equiv \sigma_C (A \times B)$$

I.e. take the cartesian product (cross product) of A and B and select tuples which satisfy condition C

© 2010 Griffith University

3

Equijoin

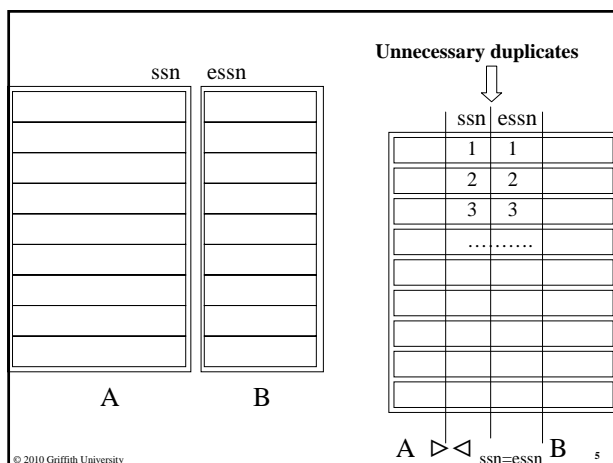
$$A \bowtie_{\text{equality conditions}} B$$

where the equality conditions prescribe that corresponding attribute values be equal

E.g. $\text{Employee} \bowtie_{\text{SSN} = \text{ESSN}} \text{Dependent}$
will join the corresponding tuples of Employee and Dependent

© 2010 Griffith University

4



Natural join

- To remove the duplicate columns a variation of equijoin is used

$$A \bowtie_{\text{equality conditions}} B$$

Example:

$$\text{Employee} \bowtie_{\text{ssn}=\text{essn}} \text{Dependent}$$

© 2010 Griffith University

6

- If no condition is listed, ***all corresponding attributes*** (ones with the same name) are joined through the equality condition
But, be careful:

Employee * Dependent

- would not work, because there are no attributes with the same name
- sometimes attribute names inadvertently join. ***It is best to always list the join condition explicitly!***

© 2010 Griffith University

7

Division

$$A \div B$$

Example:

“List the employees who work on every project”

Take a simplified company schema:

```
employee(ssn)
project(pno)
works_on(ssn,pno)
```

© 2010 Griffith University

8

- If we knew that every employee works on every project, then the works_on relation would be unnecessary, it could be reconstructed by the cartesian product
 $W = \text{employee} \times \text{project}$
- if we took away works_on from W:
 $W \setminus \text{works_on}$
the result would contain potential, but not actual works_on relationship instances

© 2010 Griffith University

9

- Therefore

$$\Pi_{\text{ssn}} (\text{Employee} \times \text{Project} \setminus \text{Works_on})$$


is the set of employees who do *not* work on every project

- To calculate the set of employees who work on every project we take this away from employee:

$$\Pi_{\text{ssn}} \text{Employee} \setminus \Pi_{\text{ssn}} (\text{Employee} \times \text{Project} \setminus \text{Works_on})$$


© 2010 Griffith University

10



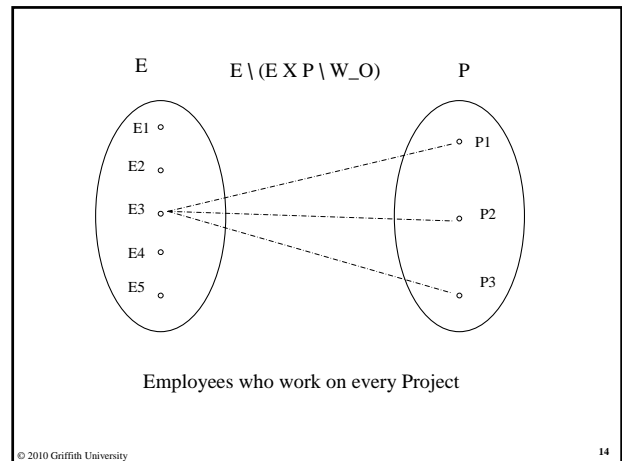
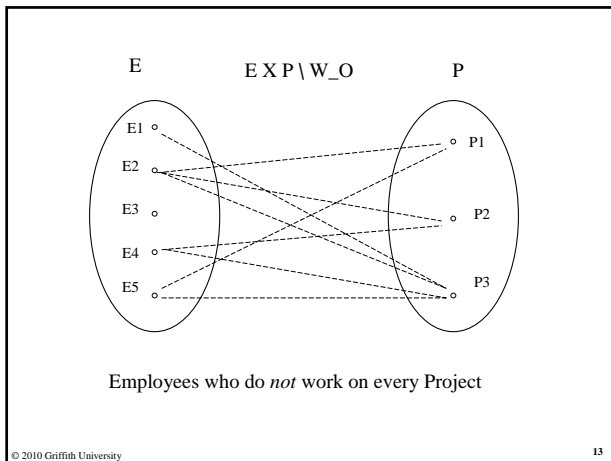
© 2010 Griffith University

11



© 2010 Griffith University

12



Introduce the division operator:

$$\text{Works_on}(\text{ssn}, \text{pno}) \div \text{Project}(\text{pno}) \equiv \Pi_{\text{ssn}} \text{Employee} \setminus (\Pi_{\text{ssn}} (\text{Employee} \times \text{Project} \setminus \text{Works_on}))$$

We can do this in general as well:

Definition

$A(a, b) \div B(b) \equiv$ the largest set $C(a)$ such that $B \times C \subseteq A$

Simplifying relational algebra expressions

- $10 * 5 + 10 * 7 + 10 * 8 = 10 * (5 + 7 + 8) = 10 * 20 = 200$
- In the same way expressions in relational algebra can be simplified based on the discussed properties of the operators
- Query processors evaluate queries using simplification or to the contrary, elaboration so that the execution is cheaper (e.g cutting down on tuples which will later be discarded anyway)

Example

Simpler

$$\begin{aligned} & \Pi_{\text{ssn}, \text{name}} (\sigma_{\text{name}='mary', \text{Employee} * \text{dno}=\text{dnumber}} \sigma_{\text{dname}='research', \text{Department}}) = \\ & \Pi_{\text{ssn}, \text{name}} (\Pi_{\text{ssn}, \text{name}, \text{dno}} \sigma_{\text{name}='mary', \text{Employee} * \text{dno}=\text{dnumber}} \Pi_{\text{dnumber}} \sigma_{\text{dname}='research', \text{Department}}) \end{aligned}$$

Cheaper to evaluate

Query plan

- Relational algebra operators are implemented as parametric programming language functions, manipulating data structures that implement sets of tuples (relations)
- A query plan is a suitable concatenation (or parallel execution model) of such functions

Query plans can be evaluated ..

- ...by the main processor of the computer on which the DBMS runs
- Simple operations, like select or project, can be performed by the disk drive itself - in order to avoid the transfer of massive amounts of data to main memory (later to be discarded anyway)

Conclusion

- Relational algebra has been introduced as an algebra of relations
- Each operator (monadic or binary) works on relations and returns relations
- Additional operations exist to carry out aggregation (this will be discussed only when studying SQL).
- Relational algebra can be used to express queries, as expressions to be evaluated on a database instance
- **Caution:** The result will depend on the actual database instance. The expression should be correct for *every / any* database instance.

The end